

Homework 2 Report

Yige Hu and Zhiting Zhu

1 P1

```
1: function REC_PSUM( $a, x_0, b, n$ )

2:   if ( $n == 1$ ) then
3:      $s(0) = x_0$ ; return; end;
4:   end if

5:    $x = \text{zeros}(n/2, 1)$ ;
6:    $a\_new = \text{zeros}(n/2 - 1, 1)$ ;
7:    $x(0) = x_0$ ;
8:   parfor  $i = 1 : n$  do
9:      $x(i) = b(i)$ ;
10:  end parfor

11:  parfor  $i = 0 : n/2 - 1$  do
12:     $y(i) = x(2 * i) * a(2 * i + 1) + x(2 * i + 1)$ ;
13:    if ( $i \neq 0$ ) then
14:       $a\_new(i) = a(2 * i) * a(2 * i + 1)$ ;
15:    end if
16:  end parfor

17:   $c = \text{REC\_PSUM}(a\_new, y(0), y[1 : n/2 - 1], n/2)$  ;

18:   $s(0) = x_0$ ;
19:  parfor  $i = 1 : n - 1$  do
20:    if isOdd( $i$ ) then
21:       $s(i) = c(i/2)$ ;
22:    else
23:       $s(i) = c((i - 1)/2) * a(i) + x(i)$ ;
24:    end if
25:  end parfor

26:  return  $s$ ;
27: end function
```

2 P2

2.1 Algorithm

```
1: function SCAN( $x, n, l$ )

2:    $step = \text{ceil}(\log_2(n))$ 
3:    $temp = n >> 1$ 
```

```

4:   offset = 1

5:   parfor i = 0 :  $n/2 - 1$  do
6:       for j = i; j < temp; j + = nthreads do
7:           indx2 = offset * (2 * i + 2) - 1
8:           indx1 = offset * (2 * i + 1) - 1
9:           x(indx2) = x(indx1) + x(indx2)
10:        end for
11:        offset* = 2
12:        temp = temp >> 1
13:    end parfor

14:    temp = 2
15:    offset >>= 1

16:    parfor i = 1 :  $n/2 - 1$  do
17:        offset >>= 1
18:        for j = i; j < temp; j + = nthreads do
19:            indx2 = offset * (2 * i + 1) - 1
20:            indx1 = offset * 2 * i - 1
21:            x(indx2) = x(indx1) + x(indx2)
22:        end for
23:        temp* = 2
24:    end parfor
25: end function

```

2.2 Result

Wall Clock Time(us)	Number of threads		
Length of Array	sequential	6 threads	12 threads
1M	15679	15500	38192
10M	156797.9	212012.5	160871
100M	730794.8	1513714	1262623.5
1B	7305516.5	14843186	12431315.5

Table 1: Wall clock execution time for different array size with different number of threads for 1D vectors

Wall Clock Time(us)	Number of threads		
Length of Array	sequential	6 threads	12 threads
1M	20525.5	79923.5	146187
10M	247284.5	539063	375131.5
100M	2046770	4615023.5	3381959

Table 2: Wall clock execution time for different array size with different number of threads for 4D vectors

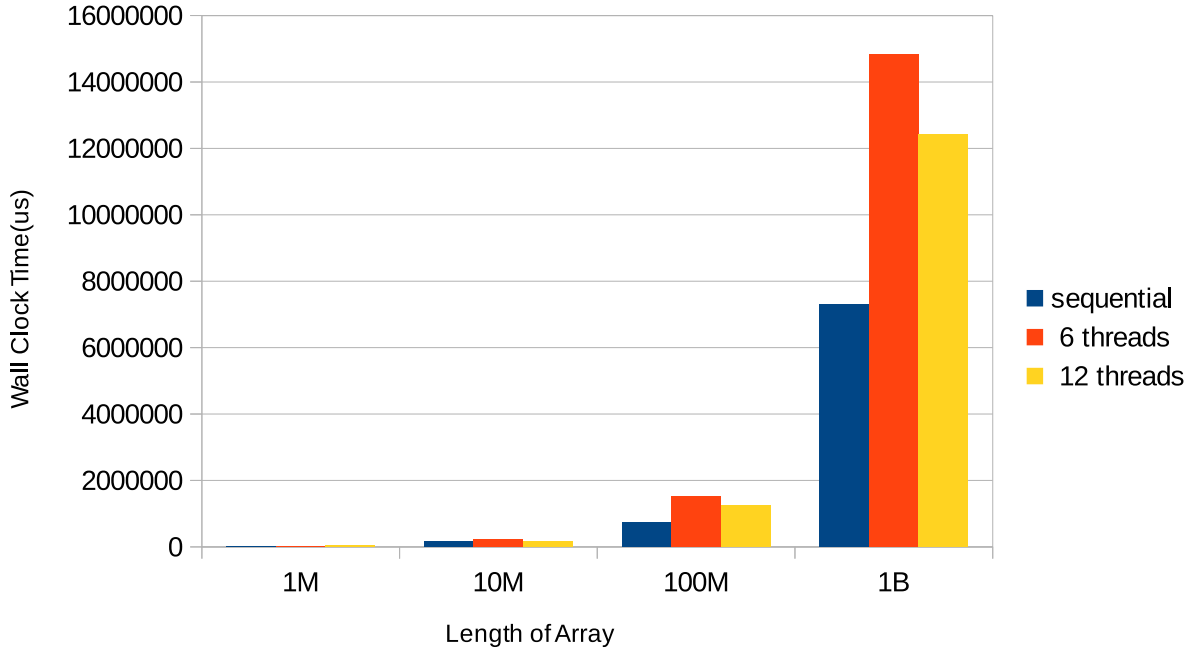


Figure 1: Wall clock execution time for different array size with different number of threads for 1D vectors

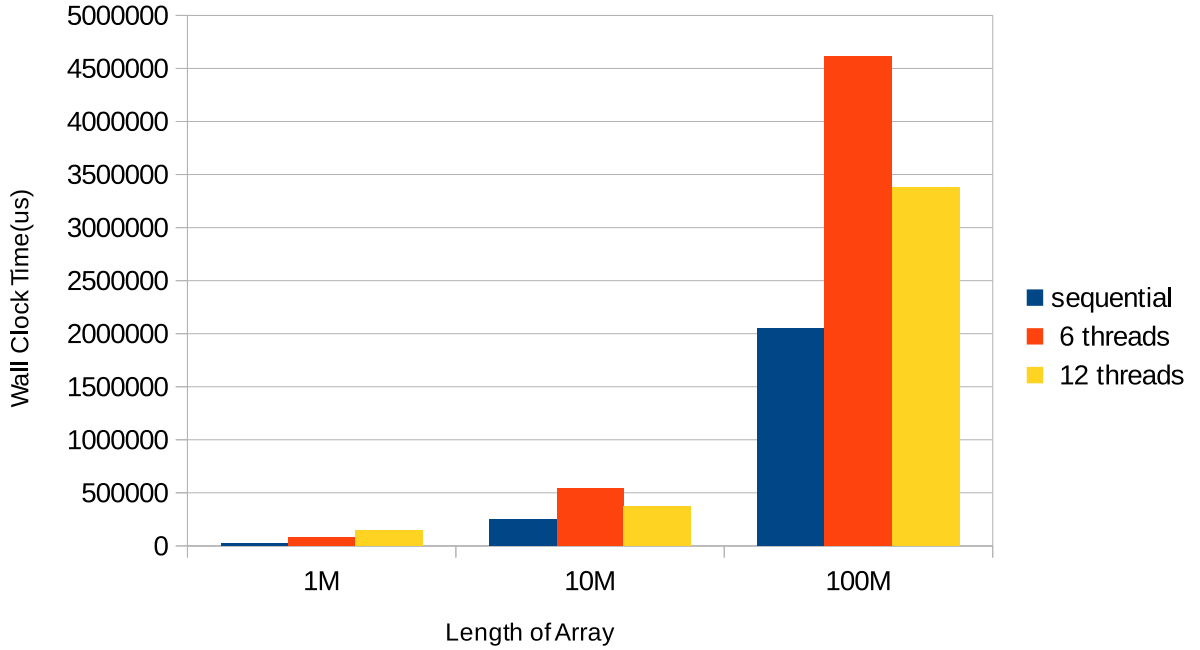


Figure 2: Timing measurements for different array size with different number of threads for 4D vectors

From Figure 1 and 2 and table 1 and 2, the parallel execution runs slower than the sequential execution. The speed up of doubling the number of threads is less than 2. The results are out of our expectation and we have consider two reasons that may be related to this result. First, the cost of administering threads

by OpenMP is high. Second, the input array is shared between multiple threads. During execution, several threads may try to write to this array at different locations which will incur frequent cache coherence for different processors as each thread is pinned to a different processor.

3 P3

3.1 Algorithm

```

1: function SEARCH(comm, keys, num_keys, sub_arr, arr_size, num_threads, rank)
2:   low_pos = arr_size * rank
3:   high_pos = low_pos + arr_size - 1
4:   low = 0
5:   high = arr_size - 1
6:   for i = 0; i < num_keys; i ++ do
7:     low = 0
8:     high = arr_size - 1
9:     position = low
10:    k = keys[i]
11:    if k < sub_arr[high] ∧ k > sub_arr[low] then
12:      if arr_size - 1 ≤ num_threads then
13:        parfor i = 1 : num_threads do
14:          if sub_arr[low + i] ≤ k then
15:            position = low
16:          end if
17:        end parfor
18:      else
19:        len = (arr_size - 1) / (num_threads + 1)
20:        position = low
21:        while len! = 1 do
22:          parfor i = 1 : num_threads do
23:            left = low + i * len
24:            right = (left + len - 1) > high ? high : (left + len - 1)
25:            if sub_arr[left] == k then
26:              position = left
27:              len = 1
28:            else if sub_arr[right] == k then
29:              position = right
30:              len = 1
31:            else
32:              if sub_arr[left] < k ∧ sub_arr[right] > k then
33:                len = (len - 1) / (num_threads + 1)
34:                low = left
35:                high = right
36:                position = low
37:              end if
38:            end if
39:          end parfor
40:        end while
41:      end if
42:      position = position + low_pos
43:      Output position
44:    else if k == sub_arr[high] then
45:      position = high + low_pos
46:      Output position

```

```

47:     else
48:         if  $k == sub\_arr[low]$  then
49:              $position = low + low\_pos$ 
50:             Output position
51:         end if
52:     end if
53: end for
54: end function

```

3.2 Result

Length	1 core	1 socket	1 node	2 node
--------	--------	----------	--------	--------

Table 3: Wall clock execution time for different configurations to search array with different length

We have successfully run the program with 16 keys but for large keyspace(2^{20} keys) our program hangs. We could not figure out the reason for hanging but we suspect that we do not have enough memory to support the full execution.